

3RCFU (Three Rivers Credit Federal Unit Data Analysis - Project Part II)

Census Neighborhood Metrics Table (First 33 Variables)

Path for the all data files required to run the code:

/anvil/projects/x-cis220051/corporate/3rivers-membership/projects/x-abajpayee

```
In [2]: import pandas as pd
import numpy as np
import os
import pickle
import shutil
from datetime import datetime, timedelta
import matplotlib.pyplot as plt
import seaborn as sns
import pyreadr
from sklearn.manifold import TSNE
from sklearn.decomposition import PCA
from sklearn.preprocessing import MinMaxScaler
from sklearn.cluster import KMeans
from tqdm import tqdm
%matplotlib inline
```

```
In [3]: census = pyreadr.read_r('Data' + '/Census_Neighborhood_Metrics_Table.rds')[N
print(census.shape)
census.head()
```

(22230, 48)

```
Out[3]:
```

	censuscode	Percent_Individual_Income_It10K	Percent_Individual_Income_10to15K	Perce
0	GKA4709381	0.171642	0.210199	
1	XPG8453176	NaN	NaN	
2	ILR6895472	0.123532	0.062636	
3	JAM2038971	0.171086	0.079555	
4	GTE6315027	0.222871	0.121566	

5 rows x 48 columns

```
In [4]: census.describe()
```

Out [4]:

	Percent_Individual_Income_lt10K	Percent_Individual_Income_10to15K	Percent_Individual_Income_15to25K
count	19215.000000	19215.000000	19215.000000
mean	0.135743	0.076392	0.051340
std	0.055477	0.035296	0.025455
min	0.000000	0.000000	0.000000
25%	0.101804	0.051340	0.025455
50%	0.127421	0.071639	0.041735
75%	0.158771	0.096345	0.067441
max	0.789583	0.380843	0.283861

8 rows × 47 columns

In [5]:

```
census.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22230 entries, 0 to 22229
Data columns (total 48 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   censuscode                                                            22230 non-null  object
1   Percent_Individual_Income_lt10K                                       19215 non-null  float64
2   Percent_Individual_Income_10to15K                                     19215 non-null  float64
3   Percent_Individual_Income_15to25K                                     19215 non-null  float64
4   Percent_Individual_Income_25to35K                                     19215 non-null  float64
5   Percent_Individual_Income_35to50K                                     19215 non-null  float64
6   Percent_Individual_Income_50to65K                                     19215 non-null  float64
7   Percent_Individual_Income_65to75K                                     19215 non-null  float64
8   Percent_Individual_Income_gte75K                                     19215 non-null  float64
9   Percent_Poverty                                                        19213 non-null  float64
10  Percent_neverMarried                                                  19215 non-null  float64
11  Percent_Married                                                       19215 non-null  float64
12  Percent_Education_HSgrad                                              19214 non-null  float64
13  Percent_Education_Somecollegeorassociate                             19214 non-null  float64
14  Percent_Education_Bachelor                                             19214 non-null  float64
15  Percent_Education_Graduateorprofessionaldegree                       19214 non-null  float64
16  Percent_FoodStamps_Household                                           19214 non-null  float64
17  Percent_GovAsst_Child_Household_SSI_SNAP_CPAI                       19123 non-null  float64
18  Percent_Unemployed                                                    19214 non-null  float64
19  Percent_Family_Poverty                                                 19210 non-null  float64
20  Percent_Medicaid                                                      19215 non-null  float64
21  Percent_HomeOwner                                                      19218 non-null  float64
22  Percent_Foreign_Born                                                   19219 non-null  float64
23  Percent_JobSector_Gov                                                  19218 non-null  float64
24  Percent_JobSector_SelfEmploy                                           19218 non-null  float64
25  Population_Density                                                     19168 non-null  float64
26  Percent_Black                                                          19167 non-null  float64
```

```

27  Percent_Native_American      19167 non-null float64
28  Percent_Asian                19167 non-null float64
29  Percent_Pacific_Islander     19167 non-null float64
30  Percent_Other                19167 non-null float64
31  Percent_gteTwoRaces          19167 non-null float64
32  Percent_Hispanic             19167 non-null float64
33  Percent_Age_lt18             19216 non-null float64
34  Percent_Age_18to24           19216 non-null float64
35  Percent_Age_gte65            19216 non-null float64
36  Income_Median                19163 non-null float64
37  Household_Income_Mean_Lowest_Quintile 18046 non-null float64
38  Household_Income_Mean_Second_Quintile 18046 non-null float64
39  Household_Income_Mean_Third_Quintile  18046 non-null float64
40  Household_Income_Mean_Fourth_Quintile 18046 non-null float64
41  Household_Income_Mean_Highest_Quintile 18046 non-null float64
42  GINI_Index                   19174 non-null float64
43  Household_Income_Median       19145 non-null float64
44  GrossRent_Median             18286 non-null float64
45  HousingUnit_Value_Median      18874 non-null float64
46  RealEstate_Taxes_Median       18833 non-null float64
47  MonthHousing_Costs_Median     19162 non-null float64
dtypes: float64(47), object(1)
memory usage: 8.1+ MB

```

Analysis to find null and missing values


```
In [6]: 100*(census.Percent_Individual_Income_lt10K.isna().sum()/census.shape[0])
```

```
Out[6]: 13.562753036437247
```

Around 13% of the values for payoffdate are missing this might be because these 28% of customers haven't paid off the loan yet. So, we can convert this feature into binary values feature, turning all missing values into '1' and the rest in '0', where '1' indicates that the loan is yet to be paid off and '0' indicates that the loan has been paid of The traditional means of imputing missing value like mean, median and mode won't really be helpful so we might not include this feature

```
In [7]: census.Percent_Individual_Income_lt10K.plot(kind='kde')
```

```
Out[7]: <AxesSubplot:ylabel='Density'>
```



```
In [8]: census.Percent_Individual_Income_lt10K.describe()
```

```
Out[8]: count      19215.000000  
mean         0.135743  
std          0.055477  
min          0.000000  
25%         0.101804  
50%         0.127421  
75%         0.158771  
max          0.789583  
Name: Percent_Individual_Income_lt10K, dtype: float64
```

Inplacing those missing values with median of the column

```
In [9]: census.Percent_Individual_Income_lt10K.fillna( census.Percent_Individual_Income_lt10K.median(), inplace=True)
census.Percent_Individual_Income_15to25K.fillna( census.Percent_Individual_Income_15to25K.median(), inplace=True)
census.Percent_Individual_Income_25to35K.fillna( census.Percent_Individual_Income_25to35K.median(), inplace=True)
census.Percent_Individual_Income_35to50K.fillna( census.Percent_Individual_Income_35to50K.median(), inplace=True)
census.Percent_Individual_Income_50to65K.fillna( census.Percent_Individual_Income_50to65K.median(), inplace=True)
census.Percent_Individual_Income_65to75K.fillna( census.Percent_Individual_Income_65to75K.median(), inplace=True)
census.Percent_Individual_Income_gte75K.fillna( census.Percent_Individual_Income_gte75K.median(), inplace=True)
census.Percent_Poverty.fillna( census.Percent_Poverty.median(), inplace=True)
census.Percent_neverMarried.fillna( census.Percent_neverMarried.median(), inplace=True)
census.Percent_Married.fillna( census.Percent_Married.median(), inplace=True)
census.Percent_Education_HSgrad.fillna( census.Percent_Education_HSgrad.median(), inplace=True)
census.Percent_Education_Somecollegeorassociate.fillna( census.Percent_Education_Somecollegeorassociate.median(), inplace=True)
census.Percent_Education_Bachelor.fillna( census.Percent_Education_Bachelor.median(), inplace=True)
census.Percent_Education_Graduateorprofessionaldegree.fillna( census.Percent_Education_Graduateorprofessionaldegree.median(), inplace=True)
census.Percent_FoodStamps_Household.fillna( census.Percent_FoodStamps_Household.median(), inplace=True)
census.Percent_GovAsst_Child_Household_SSI_SNAP_CPAI.fillna( census.Percent_GovAsst_Child_Household_SSI_SNAP_CPAI.median(), inplace=True)
census.Percent_Unemployed.fillna( census.Percent_Unemployed.median(), inplace=True)
census.Percent_Family_Poverty.fillna( census.Percent_Family_Poverty.median(), inplace=True)
census.Percent_Medicaid.fillna( census.Percent_Medicaid.median(), inplace=True)
census.Percent_HomeOwner.fillna( census.Percent_HomeOwner.median(), inplace=True)
census.Percent_Foreign_Born.fillna( census.Percent_Foreign_Born.median(), inplace=True)
census.Percent_JobSector_Gov.fillna( census.Percent_JobSector_Gov.median(), inplace=True)
census.Percent_JobSector_SelfEmploy.fillna( census.Percent_JobSector_SelfEmploy.median(), inplace=True)
census.Population_Density.fillna( census.Population_Density.median(), inplace=True)
census.Percent_Black.fillna( census.Percent_Black.median(), inplace=True )
census.Percent_Native_American.fillna( census.Percent_Native_American.median(), inplace=True )
census.Percent_Asian.fillna( census.Percent_Asian.median(), inplace=True )
census.Percent_Pacific_Islander.fillna( census.Percent_Pacific_Islander.median(), inplace=True )
census.Percent_Other.fillna( census.Percent_Other.median(), inplace=True )
census.Percent_gteTwoRaces.fillna( census.Percent_gteTwoRaces.median(), inplace=True )
census.Percent_Hispanic.fillna( census.Percent_Hispanic.median(), inplace=True )
```

Observation (Population Density):


```
census.Population_Density.describe()
```

```
In [13]: 100*(census.Population_Density.isna().sum()/census.shape[0])
```

```
Out[13]: 0.0
```

```
In [14]: census.Population_Density.plot(kind='kde')
```

```
Out[14]: <AxesSubplot:ylabel='Density'>
```



```
In [15]: census.Population_Density.describe()
```

```
Out[15]: count      22230.000000  
mean        1311.766802  
std         2963.941214  
min           0.000000  
25%         229.309446  
50%         727.320863  
75%        1384.060317  
max        70640.290323  
Name: Population_Density, dtype: float64
```

The statistical summary for the Population_Density variable in the dataset presents a distinct profile compared to the other 32 columns, which are primarily composed of percentage values related to various metrics. Unlike these percentage-based columns, the Population_Density is measured in absolute terms, providing concrete figures rather than relative percentages.

In this dataset, the Population_Density variable has a total of 22,230 entries, indicating a substantial data size. The average population density, denoted by the mean, is approximately 1,311.77. This figure represents the average number of individuals per unit area across all the measured locations. However, the standard deviation is significantly high at around 2,963.94, revealing a broad dispersion of population density values. This wide range suggests that the dataset encompasses a diverse range of areas, from sparsely populated to highly urbanized regions.

The minimum value recorded for population density is 0, which might indicate areas with negligible or no population. The 25th percentile is at 229.31, meaning a quarter of the areas have this density or lower, pointing to a number of less populated regions. The median value, or the 50th percentile, is 727.32, which is markedly lower than the mean, suggesting a right-skewed distribution. This skewness indicates that while most areas have a moderate population density, there are a few areas with extremely high densities that increase the average.

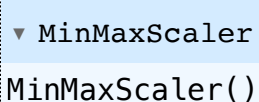
Finally, the maximum value in the dataset for population density is an exceptionally high 70,640.29, confirming the presence of highly urbanized areas within the dataset. This extreme value contrasts sharply with the overall moderate densities observed in most areas, further emphasizing the diverse range of environments covered in the dataset.

```
In [16]: normalization_features = [ 'Percent_Individual_Income_lt10K', 'Percent_Indivi
```

```
In [17]: updated_normalization_features = [ 'Percent_Individual_Income_lt10K', 'Perce
```

Normalization of the 2 lists

```
In [19]: scaler = MinMaxScaler()  
scaler.fit(census[normalization_features])  
scaler.fit(census[updated_normalization_features])
```

```
Out[19]:    
MinMaxScaler()
```


```
In [25]: normalized_census = census[normalization_features]
         updated_normalized_census = census[updated_normalization_features]
```

Determining the optimal number of clusters using elbow method

```
In [26]: inertia = []
         for n in tqdm(range(2 , 30)):
             algorithm = (KMeans(n_clusters = n , init='k-means++', n_init = 3 , max_it
                               tol=0.0001, random_state= 111 , algorithm='elkan'))
             algorithm.fit(normalized_census)
             inertia.append(algorithm.inertia_)

         plt.figure(1 , figsize = (15 ,6))
         plt.plot(np.arange(1 , 29) , inertia , 'o')
         plt.plot(np.arange(1 , 29) , inertia , '-' , alpha = 0.5)
         plt.xlabel('Number of Clusters') , plt.ylabel('Inertia')
         plt.show()
```

100% | ██████████ | 28/28 [00:18<00:00, 1.48it/s]




Clustering plots:


```
In [23]: algorithm = (KMeans(n_clusters = 4 ,init='k-means++', n_init = 3 ,max_iter=5
                             tol=0.0001, random_state= 111 , algorithm='elkan') )
algorithm.fit(normalized_census)
labels1 = algorithm.labels_
centroids1 = algorithm.cluster_centers_

pca = PCA(n_components=2)
pca.fit(normalized_census)
reduced_data = pca.transform(normalized_census)
plt.figure(1, figsize=(15,12))
visual_df = pd.DataFrame( {'x' : reduced_data[:,0], 'y' : reduced_data[:,1],
sns.scatterplot(data=visual_df, x='x', y='y', hue='labels', style='labels',
plt.title('Scatter plot of clusters with number of clusters = 4')
```

Out[23]: Text(0.5, 1.0, 'Scatter plot of clusters with number of clusters = 4')




Determining the optimal number of clusters using elbow method excluding the column, "Population Density"

```
In [22]: inertia = []
for n in tqdm(range(2, 30)):
    algorithm = (KMeans(n_clusters = n, init='k-means++', n_init = 3, max_iter=1000,
                        tol=0.0001, random_state= 111, algorithm='elkan'))
    algorithm.fit(updated_normalized_census)
    inertia.append(algorithm.inertia_)

plt.figure(1, figsize = (15, 6))
plt.plot(np.arange(1, 29), inertia, 'o')
plt.plot(np.arange(1, 29), inertia, '-', alpha = 0.5)
plt.xlabel('Number of Clusters'), plt.ylabel('Inertia')
plt.show()
```

100% | ██████████ | 28/28 [00:26<00:00, 1.05it/s]




Clustering plots:

```
In [24]: algorithm = (KMeans(n_clusters = 4, init='k-means++', n_init = 3, max_iter=500,
                            tol=0.0001, random_state= 111, algorithm='elkan'))
algorithm.fit(updated_normalized_census)
labels2 = algorithm.labels_
centroids2 = algorithm.cluster_centers_

pca = PCA(n_components=2)
pca.fit(updated_normalized_census)
reduced_data = pca.transform(updated_normalized_census)
plt.figure(1, figsize=(15, 12))
visual_df = pd.DataFrame({'x': reduced_data[:, 0], 'y': reduced_data[:, 1]})
sns.scatterplot(data=visual_df, x='x', y='y', hue='labels', style='labels',
               plt.title('Scatter plot of clusters with number of clusters = 4'))
```

Out[24]: Text(0.5, 1.0, 'Scatter plot of clusters with number of clusters = 4')



Saving the data and corresponding cluster labels

```
In [76]: solution_df = pd.DataFrame(  
    {  
        'censuscode' : census['censuscode'],  
        'labels' : labels2  
    }  
)  
solution_df.to_csv("census_labelled.csv", index=None)  
np.unique(solution_df.labels)
```

```
Out[76]: array([0, 1, 2, 3], dtype=int32)
```

Summary Report on K-Means++ Clustering Analysis of Census Data

Dataset Overview: The dataset comprises 22,230 entries, each representing a unique census code. It contains 33 columns, covering a wide range of socio-economic and

demographic indicators, such as individual income brackets, marital status, educational attainment, employment status, racial demographics, and more. Clustering Analysis:

Objective:

To segment the dataset into four distinct clusters using the k-means++ algorithm.

Methodology:

With 'Population Density': The initial clustering included all columns, notably 'Population Density'. Without 'Population Density': A second clustering analysis was conducted after removing the 'Population Density' column. Findings: Clusters with 'Population Density': The inclusion of 'Population Density' in the analysis led to clusters where this variable significantly influenced the grouping. This might have overshadowed other socio-economic factors due to the high variance or distinct patterns in population density across regions. Clusters without 'Population Density': Removing 'Population Density' resulted in clusters more reflective of socio-economic and demographic similarities, not overshadowed by geographical density factors. This approach likely offered a clearer view of how other variables, such as income levels, educational attainment, and racial composition, interact and cluster independently of geographical density considerations. Implications and Recommendations:

Comparative Insights:

The comparison between the two clustering results highlights how a single variable, like 'Population Density', can significantly skew or influence cluster formation. Decision for Analysis: The clusters formed without 'Population Density' appear to be more informative for socio-economic and demographic analysis, suggesting its exclusion is beneficial for certain types of analysis. Future Studies: It's recommended to consider the specific research goals when deciding whether to include or exclude variables like 'Population Density'. For studies focusing on urban-rural divides or regional planning, including it might be more relevant.

Conclusion:

The k-means++ clustering analysis, with and without the 'Population Density' variable, provides valuable insights into how different variables influence data segmentation. The findings underscore the importance of careful variable selection in clustering algorithms to yield the most meaningful and relevant results for the intended analysis.